

# Gradient\_Descent\_and\_Back\_Propagation.R

kriti

Sat Nov 03 21:49:54 2018

```
###understanding cost function and back propogation of errors
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
##we will run 100 iterations to update the weights
```

```
##the weight and bias in a neural network can be modelled like the slope and intercept in a linear regression
```

```
##just like in a linear regression we try to update the beta values after each iteration, in a neural network
```

```
##we update the weights in a similar way
```

```
n=100
```


```
intercept=3
```

```
slope=4.5
```


```
x=rnorm(n)
```

```
hist(x)
```


## Histogram of x



```
y=x*slope+intercept  
plot(x=x,y=y)
```



```
##adding noise  
y=x*slope+intercept+rnorm(100)  
plot(x=x,y=y)
```



```
##we add a columns of 1s to facilitate matrix mul
```

```
## its imp to notice that we want one out for each row
# of data,so essentially the row of matrix one and column of
# matrix 2 should be one.Hence since col of mat1=row of mat2
# for matrix mul hence we add another row of 1's in the data
x_b<-cbind(x,rep(1,n))
head(x_b)
```

```
##           x
## [1,] -1.4968709 1
## [2,]  2.0355967 1
## [3,]  1.1461531 1
## [4,]  0.3039515 1
## [5,]  0.4949068 1
## [6,] -0.3958989 1
```

```
##creating a dataframe
data<-data.frame(x=x_b,y=y)
data
```

##	x.x	x.V2	y
## 1	-1.496870888	1	-4.03906278
## 2	2.035596737	1	14.38659355
## 3	1.146153104	1	9.03047109
## 4	0.303951495	1	4.89377923
## 5	0.494906755	1	5.08713731
## 6	-0.395898862	1	1.32734895
## 7	0.375597595	1	5.35595300
## 8	-0.577810016	1	0.31318230
## 9	0.893397990	1	6.76529397
## 10	-0.380008023	1	1.39871745
## 11	-0.357027710	1	1.34409871
## 12	0.009639976	1	2.25811867
## 13	1.540456189	1	10.04739810
## 14	-0.797293121	1	-1.78503446
## 15	-0.624245126	1	-0.40695245
## 16	-0.181903360	1	2.88909865
## 17	0.949147924	1	6.69302747
## 18	-0.366860714	1	0.09544085
## 19	1.288300348	1	8.60190244
## 20	-0.073802780	1	1.97212918
## 21	-1.634027117	1	-4.14141346
## 22	-0.123427222	1	3.70881671
## 23	0.364782093	1	6.21597170
## 24	-0.966146811	1	-1.02145629
## 25	-1.290545651	1	-2.89493794
## 26	-1.514009960	1	-3.31883608
## 27	0.542343311	1	5.22220973
## 28	1.255756488	1	9.31507520
## 29	0.460602619	1	3.75431205
## 30	-1.731658593	1	-4.06907539
## 31	-0.166958647	1	3.63296924
## 32	-1.366072992	1	-2.03163279
## 33	0.872536580	1	6.69154365
## 34	1.135311334	1	9.64666226
## 35	-0.041428738	1	0.73279572
## 36	-0.275325909	1	2.99167283
## 37	1.589155117	1	10.24174535
## 38	-0.223390365	1	0.65752024
## 39	0.707787567	1	5.13729319
## 40	-0.566538379	1	0.63381017
## 41	1.357883868	1	10.24890141
## 42	1.305104745	1	8.67446355
## 43	-1.092030224	1	-2.50333290
## 44	-0.457279757	1	0.56224997
## 45	-0.050783085	1	3.92656043
## 46	-0.483100104	1	1.70250676
## 47	-1.923479947	1	-3.72216254
## 48	-0.876061335	1	-0.42632108
## 49	-0.314151333	1	-0.22284664
## 50	0.384494847	1	3.49189467
## 51	0.630399668	1	4.75763437
## 52	0.076010719	1	2.29805845
## 53	0.783954787	1	4.17602028
## 54	-0.442257709	1	-0.76016634
## 55	1.130085740	1	6.81460930

```
## 56 0.617862224 1 5.44210419
## 57 -0.702086902 1 1.42097336
## 58 0.804946905 1 7.47715744
## 59 0.523758280 1 5.62592500
## 60 0.586711487 1 5.69205661
## 61 -0.461080382 1 2.30385968
## 62 0.361594005 1 4.80782737
## 63 -1.079538302 1 -0.26227779
## 64 -0.977102188 1 -1.90608314
## 65 0.294670124 1 4.30821877
## 66 0.112992073 1 3.57552873
## 67 0.511993941 1 5.28513768
## 68 1.359297179 1 10.79447545
## 69 0.395124629 1 4.59389869
## 70 -0.890957517 1 -0.29844304
## 71 0.602680848 1 6.59883674
## 72 0.070465483 1 2.91141830
## 73 -0.401519797 1 0.03851908
## 74 -0.929177244 1 -0.65067484
## 75 -0.831373558 1 -2.21721978
## 76 0.805025831 1 4.20709408
## 77 0.521993733 1 4.71231766
## 78 0.538438833 1 4.29498775
## 79 -0.030217794 1 1.80130241
## 80 -0.680781130 1 -1.80711180
## 81 -0.215103177 1 1.89448539
## 82 -0.544381272 1 1.99181858
## 83 0.160722764 1 3.27471360
## 84 1.824658382 1 11.08406509
## 85 -0.037768465 1 2.26001572
## 86 -0.196833948 1 1.06068617
## 87 0.342446225 1 4.85032391
## 88 0.356618960 1 4.18685231
## 89 -0.744108798 1 -0.49091736
## 90 1.693654668 1 11.63947667
## 91 0.673389450 1 5.93178738
## 92 -0.376670160 1 0.80059107
## 93 -1.710998906 1 -5.37228665
## 94 0.376146187 1 2.88479574
## 95 -0.222982160 1 3.33560950
## 96 -0.150352581 1 2.85792469
## 97 -0.238628940 1 0.77854094
## 98 0.755413855 1 7.86769189
## 99 1.703566825 1 11.51056699
## 100 -0.595217149 1 0.13009541
```

```
##defining the Learning rate and no of iterations

learning_rate=0.05
n_iterations=100

##though the value of theta zero and theta one is 3 and 4.5
## we start the model with very high values which are not
## at all close to the original values
theta=matrix(c(20,20))
theta
```

```
##      [,1]
## [1,]  20
## [2,]  20
```

```
##creating vectors that store the values of the parameters and
## the errors

b0<-vector("numeric",length=n_iterations)
b1<-vector("numeric",length=n_iterations)
sse<-vector("numeric",length=n_iterations)

res<-x_b%%theta-y

grad=(t(x_b)%%res)*2/n
grad
```

```
##      [,1]
## x 24.85373
## 35.57653
```

```
grad*learning_rate
```

```
##      [,1]
## x 1.242686
## 1.778827
```

```
theta=theta-grad*learning_rate
theta
```

```
##      [,1]
## x 18.75731
## 18.22117
```

```
ss=(y-(x_b%%theta))**2
head(ss)
```

```
##           [,1]
## [1,]  33.83796
## [2,] 1765.42039
## [3,]  941.84267
## [4,]  362.09171
## [5,]  502.52894
## [6,]   89.63972
```

*###we see that after one iteration the value of weights have come down*

*##This we will repeat for N iterations inside a loop*

*##reinitializing theta*

```
theta=matrix(c(20,20))
theta
```

```
##           [,1]
## [1,]    20
## [2,]    20
```

```
for(i in 1:n_iterations)
{
  res<-(x_b%%theta)-y
  gradient<-(t(x_b)%%res)*2/n
  theta<-theta-(learning_rate*gradient)
  sse[i]<-sum((y-(x_b%%theta))**2)
  b0[i]<-theta[2]
  b1[i]<-theta[1]
}

model_i<-data.frame(model_iter=1:n_iterations,sse=sse,b0=b0,b1=b1)
model_i
```



##	model_iter	sse	b0	b1
## 1	1	40679.1139	18.221173	18.757314
## 2	2	33357.6724	16.626219	17.616424
## 3	3	27374.3581	15.196260	16.568808
## 4	4	22482.0113	13.914347	15.606673
## 5	5	18479.5121	12.765262	14.722894
## 6	6	15203.1470	11.735346	13.910951
## 7	7	12519.6023	10.812336	13.164883
## 8	8	10320.2780	9.985223	12.479231
## 9	9	8516.6707	9.244126	11.849001
## 10	10	7036.6244	8.580176	11.269621
## 11	11	5821.2858	7.985415	10.736905
## 12	12	4822.6294	7.452697	10.247018
## 13	13	4001.4475	6.975612	9.796449
## 14	14	3325.7150	6.548408	9.381980
## 15	15	2769.2604	6.165922	9.000662
## 16	16	2310.6847	5.823523	8.649792
## 17	17	1932.4812	5.517055	8.326895
## 18	18	1620.3198	5.242790	8.029697
## 19	19	1362.4631	4.997384	7.756116
## 20	20	1149.2922	4.777838	7.504241
## 21	21	972.9188	4.581460	7.272319
## 22	22	826.8700	4.405838	7.058742
## 23	23	705.8306	4.248808	6.862033
## 24	24	605.4332	4.108429	6.680838
## 25	25	522.0867	3.982961	6.513912
## 26	26	452.8359	3.870844	6.360113
## 27	27	395.2475	3.770681	6.218392
## 28	28	347.3163	3.681217	6.087785
## 29	29	307.3884	3.601329	5.967407
## 30	30	274.0988	3.530010	5.856444
## 31	31	246.3201	3.466358	5.754150
## 32	32	223.1202	3.409564	5.659835
## 33	33	203.7279	3.358904	5.572870
## 34	34	187.5047	3.313730	5.492673
## 35	35	173.9214	3.273459	5.418710
## 36	36	162.5390	3.237572	5.350489
## 37	37	152.9932	3.205603	5.287559
## 38	38	144.9813	3.177133	5.229504
## 39	39	138.2514	3.151791	5.175941
## 40	40	132.5940	3.129241	5.126519
## 41	41	127.8347	3.109184	5.080912
## 42	42	123.8278	3.091353	5.038824
## 43	43	120.4520	3.075507	4.999979
## 44	44	117.6059	3.061434	4.964124
## 45	45	115.2046	3.048941	4.931027
## 46	46	113.1774	3.037856	4.900473
## 47	47	111.4647	3.028027	4.872264
## 48	48	110.0170	3.019318	4.846219
## 49	49	108.7924	3.011604	4.822170
## 50	50	107.7559	3.004778	4.799961
## 51	51	106.8782	2.998742	4.779452
## 52	52	106.1344	2.993408	4.760511
## 53	53	105.5039	2.988699	4.743016
## 54	54	104.9691	2.984545	4.726856
## 55	55	104.5152	2.980884	4.711929


## 56	56	104.1298	2.977661	4.698138
## 57	57	103.8025	2.974827	4.685398
## 58	58	103.5243	2.972338	4.673628
## 59	59	103.2878	2.970155	4.662752
## 60	60	103.0866	2.968242	4.652703
## 61	61	102.9154	2.966569	4.643417
## 62	62	102.7698	2.965108	4.634835
## 63	63	102.6457	2.963835	4.626904
## 64	64	102.5400	2.962727	4.619575
## 65	65	102.4500	2.961765	4.612800
## 66	66	102.3732	2.960932	4.606539
## 67	67	102.3078	2.960213	4.600751
## 68	68	102.2520	2.959593	4.595400
## 69	69	102.2044	2.959061	4.590455
## 70	70	102.1638	2.958606	4.585882
## 71	71	102.1291	2.958219	4.581655
## 72	72	102.0995	2.957891	4.577747
## 73	73	102.0743	2.957614	4.574134
## 74	74	102.0527	2.957383	4.570793
## 75	75	102.0342	2.957190	4.567703
## 76	76	102.0185	2.957032	4.564847
## 77	77	102.0050	2.956904	4.562205
## 78	78	101.9935	2.956801	4.559763
## 79	79	101.9837	2.956720	4.557504
## 80	80	101.9753	2.956658	4.555414
## 81	81	101.9681	2.956612	4.553482
## 82	82	101.9620	2.956581	4.551695
## 83	83	101.9567	2.956561	4.550043
## 84	84	101.9522	2.956551	4.548514
## 85	85	101.9483	2.956549	4.547100
## 86	86	101.9451	2.956554	4.545792
## 87	87	101.9422	2.956565	4.544582
## 88	88	101.9398	2.956581	4.543463
## 89	89	101.9378	2.956601	4.542428
## 90	90	101.9360	2.956624	4.541470
## 91	91	101.9345	2.956649	4.540584
## 92	92	101.9332	2.956675	4.539765
## 93	93	101.9321	2.956703	4.539007
## 94	94	101.9311	2.956732	4.538305
## 95	95	101.9303	2.956762	4.537656
## 96	96	101.9296	2.956791	4.537056
## 97	97	101.9290	2.956821	4.536500
## 98	98	101.9285	2.956850	4.535986
## 99	99	101.9281	2.956879	4.535511
## 100	100	101.9277	2.956907	4.535070

```
###plotting all the lines and the final regression line
```

```
p1 <- data %>%  
  ggplot(aes(x=x, y=y)) +  
  geom_abline(aes(intercept = b0,  
                  slope = b1,  
                  colour = -sse,  
                  frame = model_iter),  
             data = model_i,  
             alpha = .50  
  ) +  
  geom_point(alpha = 0.4) +  
  geom_abline(aes(intercept = b0,  
                  slope = b1),  
             data = model_i[100, ],  
             alpha = 0.5,  
             size = 2,  
             colour = "dodger blue") +  
  geom_abline(aes(intercept = b0,  
                  slope = b1),  
             data = model_i[1, ],  
             colour = "red",  
             alpha = 0.5,  
             size = 2) +  
  scale_color_continuous(low = "red", high = "grey") +  
  guides(colour = FALSE) +  
  theme_minimal()
```

```
## Warning: Ignoring unknown aesthetics: frame
```

```
p1
```



*##plotting the decreasing sse errors*

```
p2 <- model_i[1:30,] %>%
  ggplot(aes(model_iter, sse, colour = -sse)) +
  geom_point(alpha = 0.4) +
  theme_minimal() +
  labs(x = "Model iteration",
       y = "Sum of Squared errors") +
  scale_color_continuous(low = "red", high = "dodger blue") +
  guides(colour = FALSE)
p2
```

